


A Full Case Study using SMV

CS 427I
Abhik Roychoudhury
National University of Singapore


Abhik Roychoudhury, CS427I lectures 1



So far ...

- **Basics of modeling**
 - Includes details of SMV syntax
- **Toy examples**
 - ABP, Traffic Light Controller
- We need to model/verify a medium sized problem completely to get the feel


Abhik Roychoudhury, CS427I lectures 2



Case Studies

- Many well-publicized successes of Model Checking in the verif. of processors / cache coherence protocols
 - Encore Gigamax Cache coherence protocol
 - IEEE FutureBus+ standard
 - T9000 virtual channel processor
- Our case study is a slightly more software centric coherence protocol, one for distributed file systems


Abhik Roychoudhury, CS427I lectures 3



Note of Caution

- Following slides use CMU SMV syntax which is a bit different from Cadence SMV syntax.
 - This has been done to ensure uniformity with the reading material.
 - The syntax differences are however minimal e.g.
 - MODULE x(...) instead of MODULE x(...)
 - VAR {


Abhik Roychoudhury, CS427I lectures 4



Before starting ...

- SMV employs **symbolic** model checking
 - More space efficient than the explicit state MC algorithm which proceeds by graph search.
 - Uses a data structure called Binary Decision Diagrams for compact internal representation of state space.
 - This will be covered in later lectures, but you do not need to understand symbolic MC for modeling and verifying using SMV.

Abhik Roychoudhury, CS427I lectures 5



Success of Model Checking

- Primarily in hardware verification
- Routinely used for processor verification (or modules of it) in Intel etc
- Employed by CAD giants and processor design companies
 - Cadence, Intel, Motorola ...

Abhik Roychoudhury, CS427I lectures 6

Verification or Bug Hunting

- Model Checking verifies only the design, not the implementation
- So, MC is more of a bug detection mechanism
 - But “Automated” bug detection
- This line will be more acceptable to practitioners than presenting it as a formal verification technique.

Abhik Roychoudhury, CS4271 lectures 7

Reference

- A Case Study in Model Checking Software Systems
 - Jeanette Wing and M. Vaziri-Farhana
 - FSE 1995, CMU Tech report CMU-CS-96-124
 - Another version in “Science of Computer Programming”, v20, 1997, 273-299.

Abhik Roychoudhury, CS4271 lectures 8

File System Cache coherence

- Several clients and servers
- Each file is authorized by a single server
- Clients cache files on obtaining them from the server
- Clients talk directly only to the server, not to other clients.


Abhik Roychoudhury, CS4271 lectures 9

The Cache Coherence Problem

- A client can have a copy of a file, but
 - It may not be sure that it is the latest copy
 - It knows that it is an invalid copy
- The client will
 - Request for a validation from the server
 - Request a fresh copy

Abhik Roychoudhury, CS4271 lectures 10

Client's belief about its copy



Client receives val or inval from server

Abhik Roychoudhury, CS4271 lectures 11

Client – State Variables

- **out** – request to the server.
 - This can be a request to
 - Fetch a file copy
 - Validate a file copy
- **belief**
 - the current belief of the client about the status of its cached copy.

Abhik Roychoudhury, CS4271 lectures 12

Simplification # 1

- Do not model contents of shared files.
 - The file contents does not affect the coherence mechanism.

Abhik Roychoudhury, CS4271 lectures 13

Client – State Variables

- Possible values of the “belief” variable
 - Valid
 - the client knows that the cached copy is valid
 - Invalid
 - the client knows that the cached copy is invalid
 - Suspect
 - the client is not sure of the status of the cached copy
 - Nofile
 - client does not have a cached copy

Abhik Roychoudhury, CS4271 lectures 14

Simplification # 2

- In reality
 - The client has a status for every shared file
- We model
 - Only one shared file, since the coherence issues of each file is independent
 - Hence modeling one server is enough.

Abhik Roychoudhury, CS4271 lectures 15

Client SMV description

- `MODULE client(input)`
- `VAR`
 - `out: {0, fetch, validate};`
 - `belief: {valid, invalid, suspect, nofile};`
- `ASSIGN`
 - ...
- Receives `input` from server
- `out` and `belief` are state variables

Abhik Roychoudhury, CS4271 lectures 16

Client- Trans. Rel. in SMV

- `ASSIGN`
 - `init(out) := 0;`
 - `next(out) := case`
 - `(belief = nofile) : fetch;`
 - `(belief = invalid) : fetch;`
 - `(belief = suspect) : validate;`
 - `1: 0;`
 - `esac`
- ...


Abhik Roychoudhury, CS4271 lectures 17

Client – Trans. Rel. in SMV

- `init(belief) := nofile, suspect;`
- `next(belief) :=`
 - `case`
 - `(belief=nofile)&(input=val): valid;`
 - `(belief=suspect)&(input=val): valid;`
 - `(belief=suspect)&(input=invalid): invalid;`
 - `(belief=invalid)&(input=val): valid;`
 - `1 : belief`
 - `esac`
- `belief` is updated based on current value, and input from server.

Abhik Roychoudhury, CS4271 lectures 18

Server – State Diagram



Server receives fetch, validate from Client.

Abhik Roychoudhury, CS4271 lectures 19

Server – State Variables

- belief : the belief of the server about the status of a cached copy in a particular client
 - Valid
 - Invalid
 - None : Server has no knowledge (the client might or might not have a cached copy)
- For each client i , the server should model a belief[i]
 - But we will simplify further and model only one client !!

Abhik Roychoudhury, CS4271 lectures 20

Server – State Variables

- out : the output of the server to the client
 - val: indicates to the client that its cached copy is valid
 - inval : indicates to the client that its cached copy is invalid
 - 0 : default output. Ignored by client.

Abhik Roychoudhury, CS4271 lectures 21

Simplification # 3

- Server has another state variable **valid_file**
 - Boolean variable
- Models updates to the server by other clients
 - Which need not be modeled explicitly
- Suppose C.belief = suspect
 - Then C sends a validate message to Server
- If Server has received a file update message by another client C' by now
 - Then server deems cached copy of C as invalid
 - Else ...

Abhik Roychoudhury, CS4271 lectures 22

Modeling other clients

- Whether Server receives a new update from another client is modeled by the variable **valid_file**
 - Set non-deterministically since we do not model the other clients explicitly
- Modeled one client, and exploited the star topology to implicitly model the effects of other clients on the server
 - Drastically cuts down the state space.

Abhik Roychoudhury, CS4271 lectures 23

Simplification #4 (uncommon)

- Our modeling considers only finite traces
 - “Final” states in Client and Server
 - Restricted modeling of a single session
 - But can be used for verifying invariant properties
 - $AG(\text{boolean formula of atomic propositions})$
 - If we still find a violation of such properties they would have occurred in the generic modeling with infinite execution traces as well
 - The finite execution traces modeled are possible prefixes of the actual traces if the system was modeled in details

Abhik Roychoudhury, CS4271 lectures 24

System description in SMV

- `MODULE main`
 - `VAR`
 - `Client : client(Server.out);`
 - `Server: server(Client.out);`
-
- Client module already presented
 - Let us look at the Server.

Server : SMV description

- MODULE server(input)
- VAR
 - out : {0, val, inval};
 - belief : {none, valid, invalid};
 - valid_file : boolean;
- ASSIGN
 - valid_file := {0, 1};
 - ...

Server's Transition Relation

```

• init(belief) := none;
• next(belief) :=
  case
  • (belief=none) & (input=fetch): valid;
  • (belief=none) & (input=validate) & !valid_file; valid;
  • (belief=none) & (input=validate) & !valid_file:
    _____invalid;
  • (belief=invalid) & (input=fetch): valid;
  • 1 : belief
• esac

```

Server's Transition Relation

```

• init(out) := 0;
• next(out):=
• case
•   (belief=None)&(input=fetch): val;
•   (belief=None)&(input=validate)&valid_file:=
•                                     val;
•   (belief=None)&(input=validate)&
!valid_file:
•
•   inval;
•   (belief=invalid)&(input=fetch): val;
•   1 : 0
• esac

```

Properties to Verify # 1

- If client C believes cached copy of file f is valid, it does not go to server
 - In this case, server also should believe that C's cached copy of f is valid
 - AG (C.belief=valid \Rightarrow S.belief = valid)
 - Always if C.belief=valid then S.belief = valid
 - Verified to be true by CMU SMV

Property # 2

- We can also check
 - $AG(S.belief = \text{valid} \Rightarrow C.belief = \text{valid})$
 - Otherwise, client might sometimes unnecessarily go to the server for validation
 - Inefficiency : additional traffic from client to server !!
 - SMV produces a counter-example.

SMV Counter-example

- **state I.1**
 - C.out = 0, C.belief = nofile
 - S.out = 0, S.belief = none, S.valid_file = 0
- **State I.2**
 - C.out = fetch
- **State I.3**
 - S.out = val, S.belief = valid

Abhik Roychoudhury, CS4271 lectures 31

SMV counterexample

- If client C does not have a copy of a shared file (cache miss), it requests from the server via **fetch**
- Server S sends a fresh copy and updates its belief about the status of cached copy at C
- Due to the transit delay between S and C, the client still has not updated its belief, but it will do in a few steps
 - This leads to the counter-example
 - Not a cause of concern in terms of additional traffic from the client to the server.

Abhik Roychoudhury, CS4271 lectures 32

Some useful simplifications ...

- ... employed in today's case study
 - Do not model data of the data items
 - Files in this case !
 - Model only a single data item
 - Hence model a single server
 - Model only one client
 - Other clients modeled implicitly by considering their effect on the server.

Abhik Roychoudhury, CS4271 lectures 33