# Provably Correct Code Generation: A Case Study

Qian Wang , Gopal Gupta [1,2,3]

*Department of Computer Science*
*University of Texas at Dallas*
*Richardson, TX 75083, USA*

**Abstract**

Provably correct compilation is an important aspect in development of high assurance software systems. In this paper we present an approach to provably correct compilation based on *Horn logical semantics* of programming languages and partial evaluation. We also show that continuation semantics can be expressed in the Horn logical framework, and introduce Definite Clause Semantics. We illustrate our approach by developing the semantics for the SCR specification language, and using it to (automatically) generate target code in a provably correct manner.

*Key words:* Horn logic, Denotational Semantics. Compilation

## 1 Introduction

Assuring the correctness of the compilation process is an important consideration in construction of reliable software. If the compiler generates code that is not faithful to the original program code of a system, then all our efforts spent in proving the correctness of the system are futile. Proving that target code is correct w.r.t. the program source is especially important for high assurance systems, as unfaithful target code can lead to loss of life and/or property. Considerable research has been done in this area, starting from the work of McCarthy [19]. Most efforts directed at proving compiler correctness fall into three categories:

- Those that treat the compiler as just another program and use standard verification techniques to manually or semi-automatically establish its correctness. These techniques typically employ known mathematical techniques

---

such as induction proofs, axiomatic semantics, etc. They may also use theorem provers or advanced reasoning systems to semi-automate the process. The weakness of this approach is that part of the process is manual and may introduce errors. Another weakness is that each compiler developed for each language has to be separately proved correct.

- Those that *generate* the compiler automatically from the mathematical semantics of the language. Typically the semantics used is denotational. Considerable research was done in the 70s and 80s to automatically generate compilers from the semantic definition of a language. The automatically generated compilers, however, have not been used in practice due to their slowness and/or complexity of the code generated.

- Those that use program transformation systems to transform source code into target code. This approach is related to the previous one and expresses the operational semantics of the language as term rewriting rules. These term rewriting rules can be treated as a specification for a compiler, and can be proven correct. Target code is automatically obtained by applying these term-rewriting rules to the source code. The disadvantage in this approach is that specifying the compiler operationally can be quite a lengthy process. Also, the compilation time can be quite large since a term-rewriting system will be used for executing these rules.

In this paper we develop an approach based on partial evaluation and a type of semantics called *Horn logical semantics*. Our approach is similar in spirit to semantics-based approaches, however, its basis is Horn-logical semantics [5] which possesses both an operational as well as a denotational (declarative) flavor. In the Horn logical semantics approach, both the syntax and semantics of a language is specified using Horn logic statements (or pure Prolog [26]). The semantics can be viewed dually as operational or denotational. Taking an operational view, one immediately obtains an interpreter of the language $\mathcal{L}$ from the Horn-logical semantic description of the language $\mathcal{L}$. Given a program $\mathcal{P}$ written in language $\mathcal{L}$, the interpreter obtained for $\mathcal{L}$ can be used to execute the program. Moreover, given a partial evaluator for pure Prolog, the interpreter can be *partially evaluated* w.r.t. the program $\mathcal{P}$ to obtain compiled code for $\mathcal{P}$. Since the compiled code is obtained automatically via partial evaluation of the interpreter, it is faithful to the source of $\mathcal{P}$, provided the partial evaluator is correct. The correctness of the partial evaluator, however, has to be proven only once. The correctness of the code generation process for *any* language can be certified, provided the compiled code is obtained via partial evaluation.

Given that efficient execution engines have been developed for Horn Logic (pure Prolog), partial evaluation is relatively fast. Also, the declarative nature of the Horn logical semantics allows for language semantics to be rapidly obtained.

In this paper, we further develop the Horn logical semantics approach and show that continuation semantics can also be expressed in Horn logic. Moreover, we also show that in Horn logical semantics not only the syntax but also the semantics can be expressed using the definite clause grammar notation. The semantics expressed in the DCG notation allows for the store argument to be naturally hidden. We also show that continuation semantics expressed as DCGs can be partially evaluated w.r.t. a source program to obtain target code in a provably correct manner. We illustrate this in the context of the SCR (software cost reduction) method for specifying embedded real-time systems. We assume that the reader is familiar with denotational semantics, partial evaluation, logic programming, Prolog and definite clause grammars ([25,14,26] are good references for these topics).

## 2  Horn Logical Semantics

The denotational semantics of a language $\mathcal{L}$ has three components: (i) *syntax specification:* maps sentences of $\mathcal{L}$ to parse trees; it is commonly specified as a grammar in the BNF format; (ii) *semantic algebra:* represents the mathematical objects used for expressing the meaning of a program written in the language $\mathcal{L}$; these mathematical objects typically are sets or domains (partially ordered sets, lattices, etc.) along with associated operations to manipulate the elements of the sets; (iii) *valuation functions:* these are functions mapping parse trees to elements of the semantic algebras.

Traditional denotational definitions express syntax as BNF grammars, and the semantic algebras and valuation functions using $\lambda$-calculus. In Horn Logical semantics, Horn-clauses (or pure Prolog) and constraints are used instead to specify all the components of the denotational semantics of programming languages [5]. There are three major advantages of using Horn clauses and constraints for coding denotational semantics.

First, the syntax specification trivially yields an executable parser. The BNF specification of a language $\mathcal{L}$ can be quite easily transformed to a *Definite Clause Grammar* (DCG) [26]. The syntax specification written in the DCG notation serves as a parser for $\mathcal{L}$. This parser can be used to parse programs written in $\mathcal{L}$ and obtain their parse trees (or syntax trees). Thus, the syntactic BNF specification of a language is easily turned into *executable syntax* (i.e., a parser). Note that the syntax of even context sensitive languages can be specified using DCGs [6].

Second, the semantic algebra and valuation functions of $\mathcal{L}$ can also be coded in Horn-clause Logic. Since Horn-clause Logic or pure Prolog is a declarative programming notation, just like the $\lambda$-calculus, the mathematical properties of denotational semantics are preserved. Since both the syntax and semantic part of the denotational specification are expressed as logic programs, they are both executable. These syntax and semantic specifications can be loaded in a logic programming system and executed, given a program

3

written in $\mathcal{L}$. This provides us with an interpreter for the language $\mathcal{L}$. In other words, the *denotation*[4] of a program written in $\mathcal{L}$ is executable. This executable denotation can also be used for many applications, including automated generation of compiled code.

Third, non-deterministic[5] semantics can be given to a language w.r.t. resources (e.g., time, space, battery power) consumed during execution. For example, some operations in the semantic algebra may be specified in multiple ways (say in software or in hardware) with each type of specification resulting in different resource consumption. Given a program and bounds on the resources that can be consumed, only some of the many possible semantics may be viable for that program. Resource bounded partial evaluation [1] can be used to formalize resource conscious compilation (for example, energy aware compilation) [28] via Horn Logical semantics.

Horn-logical semantics can also be used for automatic verification and consistency checking [5,7]. We do not elaborate any further since we are not concerned with verification in this paper.

In [5] we show how both the syntax and semantics of a simple imperative language (a simple subset of Pascal whose grammar is shown in Figure 1) can be given in Horn Logic. The Horn logical semantics automatically yields an interpreter. Given a program $P$, the interpreter can be partially evaluated w.r.t. $P$ to obtain $P$'s compiled code.

```
Program ::= C.
C ::= C1;C2 |
      loop while B C end while |
      if B then C1 else C2 endif |
      I := E
E ::= N | Identifier | E1 + E2 |
      E1 - E2 | E1 * E2 | (E)
N ::= 0 | 1 | 2 | ...  | 9
Identifier ::= w | x | y | z
```

**Fig. 1:** BNF grammar

A program and its corresponding code generated via partial evaluation is shown below. Note that the semantics is written under the assumption that the program takes exactly two inputs (found in variables x and y) and produces exactly one output (placed in variable z).

```
w = x;
loop while w > 0
   z = z * y ;
   w = w - 1
end while.
```

```
main(X,Y,A) :-
    initialize\_store(B),
    update(a,X,B,C),
    update(b,Y,C,D),
    update(z,1,D,E),
    access(x,E,F),
    update(w,F,E,G),
    commandwhile(G,H),
```

---

[4] We refer to the denotation of a program under the Horn-logical semantics as its *Horn logical denotation.*

[5] Non-deterministic in the logic programming sense.

```
access(z,H,A).

commandwhile(A,B) :-
    access(w,A,C),
        (0<C ->
            access(z,A,D),
            access(y,A,E),
            F is D*E,
            update(z,F,A,G),
            access(w,G,H),
            I is H-1,
            update(w,I,G,J),
            commandwhile(J,B)
        ; B=A ).
```

Notice that in the program that results from partial evaluation, only a series of memory `access`, memory `update`, arithmetic and comparison operations are left, that correspond to `load`, `store`, arithmetic, and comparison operations of a machine language. The while-loop, whose meaning was expressed using recursion, will (always) partially evaluate to a *tail-recursive* program. These tail-recursive calls are easily converted to iterative structures using jumps. Though, the compiled code generated is in Prolog syntax, true machine code is only a few simple steps away. The code generation process is provably correct, since target code is obtained automatically via partial evaluation. Of course, we need to ensure that the partial evaluator works correctly. However, this needs to be done only once. Note that once we prove the correctness of the partial evaluator, compiled code for programs written in any language can be generated as long as the Horn-logical semantics of the language is given.

## 3   Definite Clause Semantics

Note that in the code generated, the `update` and `access` operations are parameterized on the memory store (i.e., they take an input store and produce an output store). Of course, real machine instructions are not parameterized on store. This problem can be solved by using the DCG notation for expressing the valuation predicates as well.

All valuation predicates take a store argument as input, modify it per the semantics of the command under consideration and produce the modified store as output [5]. Because the semantic rules are stated compositionally, the store argument "weaves" through the semantic sub-predicates called in the rule. This suggests, we can express the semantic rules in the DCG notation. Thus, we can view the semantic rules as computing the *difference* between the output and the input stores. This difference reflects the effect of the command whose semantics is being given. Expressed in the DCG notation, the store argument

is hidden away. For example, in the DCG notation the valuation predicate

```
command(comb(C1,C2),Store,Outstore):-
    command(C1,Store,Nstore),command(C2,Nstore,Outstore).
```

is written as:

```
command(comb(C1,C2)) --> command(C1), comm(C2).
```

The complete semantics of the subset of Pascal considered earlier is shown below as a DCG.

```
prog_eval(p(Comm),Vala,Valb,Output)-->update(x,Vala),
    update(y,Valb),command(Comm),access(z,Output).
command(comb(C1,C2))-->command(C1),command(C2).
command(while(B,C)) -->bool(B,Bval),
    ({Bval = true} ->
            command(C),command(while(B,C));[]).
command(ce(B,C1,C2))-->bool(B,Bval),
    ({Bval = true} -> command(C1); command(C2)).
command(assign(I,E))-->
    expression(E,Val),update(I,Val).
expression(add(E1,E2),Result)-->
    expression(E1,Val_E1),
    expression(E2, Val_E2),
    {Result is Val_E1+Val_E2}.
expression(sub(E1, E2), Result) -->
    expression(E1, Val_E1),
    expression(E2, Val_E2),
    {Result is Val_E1-Val_E2}.
expression(mul(E1, E2), Result) -->
    expression(E1, Val_E1),
    expression(E2, Val_E2),
    {Result is Val_E1*Val_E2}.
expression(id(X), Result) --> access(X, Result).
expression(num(X), X) --> [].
bool(greater(E1, E2), Bval) -->
    expression(E1, Eval1),
    expression(E2, Eval2),
    {(Eval1 > Eval2 -> Bval = true; Bval = false)}.
bool(less(E1, E2), Bval) --> expression(E1, Eval1),
    expression(E2, Eval2),
    {(Eval1 < Eval2 -> Bval = true; Bval = false)}.
bool(equal(E1, E2), Bval) --> expression(E1, Eval1),
    expression(E2, Eval2),
    {(Eval1 = Eval2 -> Bval = true; Bval = false)}.
```

Semantics expressed as a DCG

Expressed in the DCG notation, the semantic rules become more intuitively obvious. In fact, these rules have more natural reading; they can

6

be read as simple rewrite rules. Additionally, now we can partially evaluate this DCG w.r.t. an input program, and obtain compiled code that has the store argument hidden. The result of partially evaluating this DCG-formatted semantics is shown in Figure 2. Notice that the store argument weaving through the generated code shown in the original partially evaluated

```
main(X,Y,A) :-
      update(a,X),
      update(b,Y),
      update(z,1),
      access(x,F),
      update(w,F),
      commandwhile,
      access(z,A).
commandwhile :-
      (access(w,C),
      0<C ->
          access(z,D),
          access(y,E),
          F is D*E,
          update(z,F),
          access(w,H),
          I is H-1,
          update(w,I),
          commandwhile
        ; true).
```

**Fig. 2:** Store eliminated

code is gone. Notice also that the basic operations (such as comparisons, arithmetic, assignment, etc.) that appear in the target code are placed in braces in definite clause semantics, so that the two store arguments are not added during expansion to Prolog. The constructs appearing within braces can be regarded as the "terminal" symbols in this semantic evaluation, similar to terminal symbols appearing in square brackets in the syntax specification. In fact, the operations enclosed within braces are the primitive operations left in the residual target code after partial evaluation. (Note, however, that these braces can be eliminated by putting wrappers around the primitive operations; these wrappers will have two redundant store arguments that are identical, per the requirements of the DCG notation. Note also that the `-->` arrow of the DCG notation was replaced by the `:-` prior to partial evaluation.)

## 4 Continuation Semantics

So far we have modeled only direct semantics [25] using Horn logic. It is well known that direct semantics cannot naturally model exception mechanisms and `goto` statements of imperative programming languages. To express such constructs naturally, one has to resort to continuation semantics. Continuation semantics are also easily modeled in Horn Logic in the DCG format. In the definite clause continuation semantics, semantics of constructs is given in terms of the *differences of parse trees* (i.e., difference of the input parse tree and the continuation's parse tree) [27]. Each semantic predicate thus relates an individual construct (difference of two parse trees) to a fragment of the store (difference of two stores). Thus, semantic rules are of the form:

```
command(C1, C2, Program, S1, S2) :- ...
```

where the difference of C1 and C2 (say $\Delta C$) represents the command whose semantics is being given, and the difference of S1 and S2 represents the store which reflects the incremental change ($\Delta S$) brought about to the store by the command $\Delta C$. Note that the `Program` parameter is needed to carry the

mapping between labels and the corresponding command. Each semantic rule thus is a stand alone rule relating the difference of command lists, $\Delta C$, to difference of stores, $\Delta S$. *If we view a program as a sequence of commands then its semantics can simply be obtained by adding the difference of stores for each command.* That is, if we view a program $P$ as consisting of sequence of commands:

$\quad P = \Delta C_1 + \Delta C_2 + \ldots + \Delta C_n$

then its semantics $S$ is viewed as a combination of the corresponding differences of stores:

$\quad S = \Delta S_1 + \Delta S_2 + \ldots + \Delta S_n$

and the continuation semantics simply maps each $\Delta C_i$ to the corresponding $\Delta S_i$. Additionally, continuation semantics allow for cleaner, more intuitive declarative semantics for imperative constructs such as exceptions, catch/throw, goto, etc. [25]. Finally, note that the above semantic rule can also be written in the DCG notation causing the arguments S1 and S2 to become hidden:

```
command(C1, C2, Program) --> ...
```

We do not give details due to lack of space. More details can be found elsewhere [27]. However, we give below the continuation semantics of the subset of Pascal considered earlier after extending it with statement labels and a goto statement. Note that the syntax trees are now represented as a list of commands. Each command is represented in the syntax tree as a pair, whose first element is a label (possibly null) and the second element is the command itself. Only the valuation functions for commands are shown (those for expressions, etc., are similar to the one shown earlier).

```
prog([ ], _, _, 0) --> [ ].
prog(CommList,Val_x,Val_y,Output) -->
     initialize_store, update(x, Val_x), update(y, Val_y),
     comm_list(CommList,CommList), access(z,Output).


comm_list([ ],_) --> [ ].
comm_list(CurrList,Program)-->
   comm(CurrList,Rest,Program), comm_list(Rest,Program).
comm([(_,while(B,LoopBody))|T],T,Program)-->
   fix(if(B,LoopBody),Program).
comm([(_,ce(B,C1,C2))|T],T,Program) -->
   ((bool(B)) -> comm_list(C1,Program); comm_list(C2,Program)).
comm([(_,ce(B,C))|T],T,Program) -->
   ((bool(B)) -> comm_list(C,Program); [ ]).
comm([(_,goto(ID))|_],GotoCont,Program)-->
     {find_label(ID,Program,GotoCont)}.
comm([(_,assign(id(I), E))|T],T,_) --> expr(E,Val),update(I,Val).
comm([(_,abort)|_],[ ],_) --> [ ].


fix(if(B,LoopBody),Program)-->
```

8

```
((bool(B))-> comm_list(LoopBody,Program),
          fix(if(B,LoopBody),Program); [ ]).
```

The code for find_label/3 predicate is not shown. It looks for the program segment that is a target of a goto and changes the current continuation to that part of the code. Consider the program shown below to the left: The result of partially evaluating the interpreter obtained from the semantics w.r.t. this program (containing a goto) is shown below to the right.

```
z = 1;                        main(A, B, C) :-
w = x;                              initialize_store,
goto L1;                            update(x, A),
loop while w > 0                    update(y, B),
       z = z * y;                   update(z, 1),
   L1: w = w - 1                    access(x, D),
endloop while;                      update(w, D),
z = 8.                             access(w, E),
                                   F is E-1,
                                   update(w, F),
                                   fix_ifgreatidw1,
                                   update(z, 8),
                                   access(z, C).
                              fix_ifgreatidw1 :-
                                  (   access(w, A),
                                      0<A ->
                                      access(z, B),
                                      access(y, C),
                                      D is B*C,
                                      update(z, D),
                                      access(w, E),
                                      F is E-1,
                                      update(w, F),
                                      fix_ifgreatidw1
                                  ;   true
                                  ).
```

# 5  A Case Study in SCR

We have applied our approach to a number of practical applications. These include generating code for parallelizing compilers in a provably correct manner [5], generating code for controllers specified in Ada [15] and for domain specific languages [8] in a provably correct manner, and most recently generating code in a provably correct manner for the Software Cost Reduction (SCR) framework. In the rest of the paper, we show that Horn logical semantics can be practically applied; we apply it to the domain specific language of SCR,

discussed next.

The SCR (Software Cost Reduction) requirements method is a software development methodology introduced in the 80s [9] for engineering reliable software systems. The target domain for SCR is real-time embedded systems. The SCR method has been extended to describe not only functional requirements (the values the system assigns to outputs) but also nonfunctional (e.g., timing and accuracy) requirements. A number of automatic tools have been developed to aid in formal specification, verification and validation of systems using the SCR method [12]. SCR has been applied to a number of practical systems, including avionics system (the A-7 Operational flight Program), a submarine communication system, and safety-critical components of a nuclear power plant [12]. The SCR method is scalable and has been applied to document requirements of the Lockheed's C-130J Operational Flight Program which resulted in approximately 100,000 lines of Ada code.

The SCR method describes system behavior by a mathematical relation between monitored variables and controlled variables. This relation is concisely specified using *condition*, *events* and *tables*. A condition is a predicate defined on one or more variables in the specification. An event occurs when any variable changes values. The environment changes monitored values and causes input events. In response, the system updates the value of one or more controlled variable according to some relations. Each SCR table specifies the required value of a variable as a mathematical function defined on conditions and events. There are three kinds of table used in SCR: *condition* tables, *event* tables, and *mode transition* tables. The tables facilitate industrial application of the SCR method since engineers find tables relatively easy to understand and to develop [12]. In additional, tables can describe large quantities of requirement information concisely.

## 5.1 The Four-Variable Model

There are several versions of SCR. One of the most important versions is the Four Variable Model. The Four Variable Model [3], illustrated in Figure 3, represents requirements as a set of mathematical relations on four sets of variables (monitored, controlled, input, and output variables). A monitored variable represents an environmental quantity that influences system behavior, while a controlled variable represents an environmental quantity the system controls. A black box specification of required behavior is given as two relations (REQ and NAT) from the monitored quantities to the controlled quantities (rather than inputs to outputs). NAT, which defines the set of possible values, describes the natural constraints on the system behavior, such as constraints imposed by physical laws and the system environment. REQ defines the additional constraints on the system to be built as relations the system must maintain between the monitored and the controlled quantities [12].
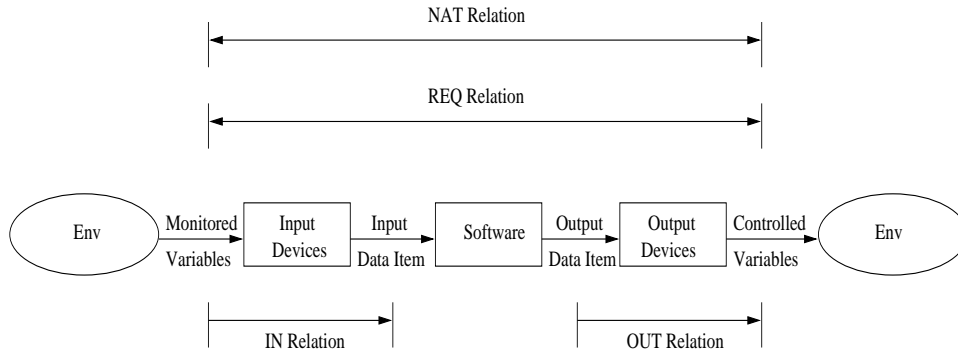
10

Fig. 3. The Four-Variable Model

A domain specific language [17] has been designed to write SCR specifications using the four variable model, as well as a large number of tools have been developed to help in checking consistency of the requirement specifications [12]. While the consistency of requirement specifications can be checked using these tools, a hurdle still remains in having absolute confidence in the final system obtained. This hurdle pertains to ensuring that the compilation process is provably correct, i.e., after consistency checking, when the SCR specification is translated into executable code then making sure that the code generated is faithful to the original specification.

We have applied our method discussed in this paper to overcome this hurdle. A Horn logical semantics for the SCR domain specific language was developed. This semantics consists of the syntax specification, semantic algebra and valuation predicates. The semantic algebra consists of operations for accessing and updating the store (values of variables as well as their type) and maintaining the various environments. Development of this semantics required just a few weeks of work (a significant part of this time was spent understanding the SCR method). The DCG notation was used both for syntax as well as semantics.

The grammar of SCR consists of five sections (type definitions, constant definitions, variable declarations, assumptions and assertions, and function definitions). User-defined data types are listed in the type definitions section. There are two types of user define data types: (i) enumerated type and (ii) integer type associated with a range. Variable declarations can include four types of variables: monitored variables, controlled variables, term variables and mode classes. The assumptions and assertions section contains predicates describing relations between variables, i.e., each assumption or assertion is a logical formula. The violation of an assumption indicates that the input does not obey the assumed environmental constraints. If an assertion is violated, it means that the specification does not satisfy a property that is was expected to satisfy. Function in SCR are defined by either a condition or an event table. Functions are used to update values of dependent variables when a monitored variable changes. The DCG for SCR has been developed in accordance to the BNF grammar supplied to us by Naval Research Lab researchers. The seman-
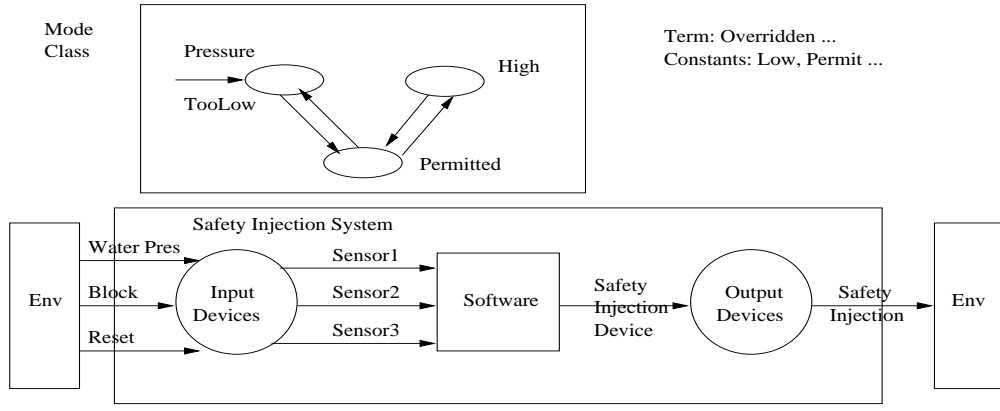
11

Fig. 4. Requirements Spec. for Safety Injection

tics of SCR's DSL is given in terms of the store semantic algebra extended with type information.

To illustrate generation of code for SCR in a provably correct manner, we consider a simplified version of the control system for safety injection described in [12]. The system uses three sensors to monitor water pressure and adds coolant to the reactor core when the pressure falls below some threshold. The system operator blocks safety injection by turning on a "block" switch and resets the system after blockage by turning on a "reset" switch. Figure 4 shows how SCR constructs could be used to specify the requirements of the control system. Water pressure and the "block" and "reset" switches are represented as monitored variables called `WaterPres`, `Block`, and `Reset`. Safety injection is represented as a controlled variable called `SafetyInjection`. Each sensor represents an input. The hardware interface between the control system software and the safety injection system serves as output. A mode class `Pressure` and a term `Overridden` help make the specification of the safety injection system concise. `Pressure` has three modes: `TooLow`, `Permitted`, and `High`. A drop in water pressure below a constant `Low` causes the system to enter mode `TooLow`; an increase in pressure above a larger constant `Permit` causes the system to enter mode `High`. The term `Overridden` is `true` if safety injection is blocked, it is `false` otherwise. An example of a condition in the specification is `WaterPres < Low`. Two examples of events are the input event `event@T(Block=on)` (the operator turns Block from `off` to `on`) and the conditioned event `@T(Block=On) WHEN WaterPres < Low` (the operator turns `Block` to on when water pressure is below `Low`). The program corresponding to this system written in the SCR domain specific language was also supplied to us by researchers at the Naval Research Labs and is shown in Appendix I. More details about SIS can be found elsewhere [12].

The Horn logical semantics developed for SCR immediately provides us with an interpreter on which the program above can be executed. Further, the interpreter was partially evaluated w.r.t. this program using the Mixtus system, and compiled code was obtained. The partially evaluated code gen-

erated that corresponds to the safety injection system is shown in Appendix II. The whole partial evaluation (using definite clause continuation semantics of SCR) required 27.1 seconds on a Sun Fire 880 with 150 MHz clock-speed and 1 CPU and 2 GB memory and generated 367 lines of assembly code in Prolog syntax (shown in Appendix II). In [17], given the same example code (Appendix I), a *relation-based* strategy (that associates C code as an attribute with parse tree nodes) required 20 minutes to generate C code, while a transformation-based method using the APTS system [21] took four minutes to generate 293 lines of C code (execution done on a SUN Ultra 450 with 2 UltraSPARC-II 296MHz CPUS and 2GB memory, running Solaris 5.6 [17]).

Even though our respective experiments have been done on different machines, the machines are comparable in speed. As can be noticed, the time taken to generate code in our case is considerably better. Note that we did not optimize the semantics at all to make it more amenable to partial evaluation as that would have reduced the readability of the semantics.

# 6 Related Work

Considerable work has been done in manually or semi-mechanically proving compilers correct. Most of these efforts are based on taking a specific compiler and showing its implementation to be correct. A number of tools (e.g., a theorem prover) may be used to semi-mechanize the proof. Example of such efforts range from McCarthy's work in 1967 [19] to more recent ones [22,11,2]. As mentioned earlier, these approaches are either manual or semi-mechanical, requiring human intervention, and therefore not completely reliable enough for engineering high-assurance systems. "Verifying Compilers" have also been considered as one of the grand challenge for computing research [13], although the emphasis here is more on developing a compiler that can verify the assertions inserted in programs (of course, such a compiler has to be proven correct first).

Considerable work has also been done on generating compilers automatically from language semantics [25]. However, because the syntax is specified as a (non-executable) BNF and semantics is specified using $\lambda$-calculus, the automatic generation process is very cumbersome. The approach outlined in this paper falls in this class, except that it uses Horn logical semantics which, we believe and experience suggests, can be manipulated more efficiently.

Considerable work has also been done in using term rewriting systems for transforming source code to target code. In fact, this approach has been applied by researchers at NRL to automatically generate C code from SCR specification using the APTS [21] program transformation system. As noted earlier, the time taken is considerably more than in our approach. Other approaches that fall in this category include the HATS system [29] that use tree rewriting to accomplish transformations. Other transformation based approaches are mentioned in [17].

13

Recently, Pnueli at al have taken the approach of verifying a given run of the compiler rather than a compiler itself [23]. This removes the burden of maintaining the compiler's correctness proof; instead each run is proved correct by establishing a refinement relationship. However, this approach is limited to very simple languages. As the authors themselves mention, their approach "seems to work in all cases that the source and target programs each consist of a repeated execution of a single loop body ..," and as such is limited. For such simple languages, we believe that a Horn logical semantics based solution will perform much better and will be far easier to develop. Development of the refinement relation is also not a trivial task. For general programs and general languages, the scalability of the approach is not established.

## 7 Conclusions

In this paper we presented an approach based on denotational semantics, Horn logic, and partial evaluation for obtaining provably correct compiled code. We illustrated our approach in the context of the SCR method for specifying real-time embedded system. The complete syntax and semantic specification for SCR was developed and used for automatically generating code for SCR specifications. Our method produces executable code considerably faster than other transformation based methods for automatically generating code for SCR specifications.

## Acknowledgement

## References

[1] S. Debray. Resource bounded partial evaluation. PEPM 1997. pp. 179-192.

[2] A. Dold, T. Gaul, W. Zimmermann Mechanized Verification of Compiler Backends Proc. Software Tools for Technology Transfer (STTT'98), Aalborg, Denmark, July 12-13, 1998.

[3] S. R. Faulk. State Determination in Hard-Embedded Systems. Ph.D. Thesis, Univ. of NC, Chapel Hill, NC, 1989.

[4] Y. Futamura. Partial Evaluation of Computer Programs: An approach to compiler-compiler. *J. Inst. Electronics and Comm. Engineers, Japan.* 1971.

[5] G. Gupta "Horn Logic Denotations and Their Applications," *The Logic Programming Paradigm: A 25 year perspective.* Springer Verlag. 1999:127-160.

14

[6] G. Gupta, H-F. Guo, A. Karshmer, E. Pontelli, et al. Semantic-Based Filtering: Logic Programming's Killer App? 4th International Symposium on Practical Aspects of Declarative Languages, LNCS 2257, Springer Verlag, pp. 82-100, Jan. 2002.

[7] G. Gupta, E. Pontelli. A Constraint-based Denotational Approach to Specification and Verification of Real-time Systems. In *Proc. IEEE Real-time Systems Symposium*, pp. 230-239. Dec. 1997.

[8] G. Gupta, E. Pontelli. A Logic Programming Framework for Specification and Implementation of Domain Specific Languages. In *Essays in Honor of Robert Kowalski*, 2003, Springer Verlag LNAI,

[9] K. L. Henninger. Specifying software requirements for complex systems: New techniques and their application. IEEE Trans. on Software Engineering. SE-5, 1. pp. 2-13.

[10] C. Gunter. Programming Language Semantics. MIT Press. 1992.

[11] J. Hannan, F. Pfenning. Compiler Verification in LF. *Proc. Seventh Annual IEEE Symposium on Logic in Computer Science.* pp. 407–418. 1992.

[12] C. L. Heitmeyer, R. D. Jeffords, and B. G. Labaw. Automated Consistency Checking of Requirements Specifications. ACM Trans. Software Eng. and Methodology 5, 3, July 1996.

[13] C. A. R. Hoare. The Verifying Compiler: A Grand Challenge for Computing Research. J.ACM, 50(1):63-69. Jan 2003.

[14] N. Jones. Introduction to Partial Evaluation. In *ACM Computing Surveys.* 28(3):480-503.

[15] L. King, G. Gupta, E. Pontelli. Verification of BART Controller: An Approach based
on Horn Logic and Denotational Semantics. In *High Integrity Software*, V. Winter and S. Bhattacharya (eds), April 2001, Kluwer Academic.

[16] M. Leuschel, B. Martens, and D. De Schreye. Controlling Generalization and Polyvariance in Partial Deduction of Normal Logic Programs. ACM Transactions on Programming Languages and Systems (TOPLAS), 20(1):208-258.

[17] E. I. Leonard and C. L. Heitmeyer. Program Synthesis from Requirements Specifications Using APTS. Kluwer Academic Publishers, 2002.

[18] K. Marriott and P. Stuckey. Constraint Programming. MIT Press, 1998.

[19] J. McCarthy and J. Painter. Correctness of a Compiler for Arithmetic Expressions. MIT AI Lab Memo, 1967.

[20] S. P. Miller. Specifying the model Logic of a Flight Guidance Systems in CoRE and SCR. Pages:44-53 Series Proceeding Article, 1998

[21] R. Paige. Viewing a Program Transformation System at Work. *Proc. Programming Language Implementation and Logic Programming*, Springer, LNCS 844. 1994.

[22] C. Pusch. Verification of Compiler Correctness for the WAM. *Proc 9th International Conference on Theorem Proving in Higher Order Logics (TPHOL'96)*. Springer-Verlag LNCS 1125, 1996.

[23] A. Pnueli, M. Siegel, E. Singerman. Translation Validation. *Proc TACAS'98*, Springer Verlag LNCS, 1998.

[24] D. Sahlin. An Automatic Partial Evaluator for Full Prolog. Ph.D. Thesis. 1994. Royal Institute of Tech., Sweden. (available at www.sics.se)

[25] D. Schmidt. *Denotational Semantics: a Methodology for Language Development.* W.C. Brown Publishers, 1986.

[26] L. Sterling & S. Shapiro. The Art of Prolog. MIT Press, '94.

[27] Q. Wang, G. Gupta. Horn Logical Continuation Semantics. UTD Technical Report. Forthcoming.

[28] Q. Wang, G. Gupta. Resource Bounded Compilation via Constrained Partial Evaluation. UTD Technical Report. Forthcoming.

[29] V. L. Winter. Program Transformation in HATS. *Software Transformation Systems Workshop*, '99.

[30] V. Winter et al. Bay Area Rapid Transit District Advance Automated Train Control System: Case Study Description. In *High Integrity Software*, V. Winter and S. Bhattacharya (eds), April 2001, Kluwer Academic.

# Appendix I: Example SCR Code

```
spec Safety_Injection_System
type definitions
  ySwitch: enum in {Off, On};
  type_mcPressure: enum in {TooLow, Permitted,High};
  yWPres: integer in [0, 2000];
constant definitions
  Low=900:integer;
  Permit=1000:integer;
monitored variables
  mWaterPres: yWPres, initially 0;
  mBlock, mReset: ySwitch, initially Off;
controlled variables
  cSafety_Injection: ySwitch, initially On;
term variables
  tOverridden: boolean, initially false;
mode classes
  mcPressure:  type_mcPressure, initially TooLow;
assumptions
   A1: (mWaterPres' >= mWaterPres AND mWaterPres'
        -  mWaterPres <=10)
   OR  (mWaterPres' < mWaterPres AND mWaterPres
        - mWaterPres' <= 10)
function definitions
var mcPressure :=
  case mcPressure
    [] TooLow
       ev
        [] @T(mWaterPres >= Low) -> Permitted
       ve
    [] Permitted
       ev
        [] @T(mWaterPres >= Permit) -> High
        [] @T(mWaterPres <  Low) -> TooLow
       ve
    [] High
       ev
        [] @T(mWaterPres < Permit) -> Permitted
       ve
  esac
var tOverridden  :=
    ev
    [] @T(mBlock=On) WHEN mReset=Off
         AND NOT ( mcPressure = High) -> true
    [] @T(mReset=On) WHEN NOT ( mcPressure = High)
         OR @T(mcPressure = High)
         OR @T(NOT (mcPressure = High) ) -> false
    ve
var cSafety_Injection ==
   case mcPressure
     [] TooLow
       if
         [] tOverridden -> Off
         [] NOT tOverridden -> On
       fi
     [] Permitted, High
       if
         [] true -> Off
         [] false -> On
       fi
      esac
```

17

# Appendix II: Generated Code

```
interpreter(A, B) :-
    interpreter1(A, B).
interpreter1(A, _) :-
    initialize_store,
    set('Low', 900),
    set(prime_Low, 900),
    set('Permit',1000),
    set(prime_Permit,1000),
    set(mWaterPres,0),
    set(prime_mWaterPres,0),
    set(mReset,'Off'),
    set(prime_mReset,'Off'),
    set(mBlock,'Off'),
    set(prime_mBlock,'Off'),
    set(cSafety_Injection,'On'),
    set(prime_
        cSafety_Injection,'On'),
    set(tOverridden, false),
    set(prime_
        tOverridden,false),
    set(mcPressure, 'TooLow'),
    set(prime_
        mcPressure,'TooLow'),
    readInputVar(A),
    access(prime_mWaterPres,B),
    access(mWaterPres,C),
    (   C=<B ->
        D=true
    ;   D=false
    ),
    access(prime_mWaterPres,E),
    access(mWaterPres, F),
    G is E-F,
    (   10<G ->
        H=false
    ;   H=true
    ),
    (   D==true,
        H==true ->
        I=true
    ;   I=false
    ),
    access(prime_mWaterPres,J),
    access(mWaterPres, K),
    (   J<K ->
        L=true
    ;   L=false
    ),
    access(mWaterPres,M),
    access(prime_
        mWaterPres,N),
    O is M-N,
    (   10<O ->
        P=false
    ;   P=true
    ),
    (   L==true,
        P==true ->
        Q=true
    ;   Q=false
```

```
    ),
    (   I==false,
        Q==false ->
        R=false
    ;   R=true
    ),
    set('A1', R),
    access(mcPressure, S),
    (   S='TooLow' ->
        access(prime_
            mWaterPres,T),
        access(prime_Low,U),
        (   U=<T ->
            V=true
        ;   V=false
        ),
        access(mWaterPres,W),
        access('Low', X),
        (   X=<W ->
            Y=true
        ;   Y=false
        ),
        (   Y==false,
            V==true ->
            Z=true
        ;   Z=false
        ),
        (   Z==true ->
            A1='Permitted'
        ;   A1=none
        )
    ;   A1=none
    ),
    (   A1==none ->
        (   S='Permitted' ->
            access(prime_
                mWaterPres,B1),
            access(prime_
                Permit,C1),
            (   C1=<B1 ->
                D1=true
            ;   D1=false
            ),
            access(mWaterPres,
                E1),
            access('Permit',F1),
            (   F1=<E1 ->
                G1=true
            ;   G1=false
            ),
            (   G1==false,
                D1==true ->
                H1=true
            ;   H1=false
            ),
            (   H1==true ->
                I1='High'
            ;   I1=none
            ),
            (   I1==none ->
                access(prime_
                    mWaterPres,J1),
                access(prime_Low,
```

18

```
            K1),                            mcPressure, Y1)
(   J1<K1 ->                    ),
    L1=true                    access(prime_mBlock,A2),
;   L1=false                   (   A2=='On' ->
),                                 B2=true
access(mWaterPres,             ;   B2=false
    M1),                       ),
access('Low',N1),              access(mBlock,C2),
(   M1<N1 ->                   (   C2=='On' ->
    O1=true                        D2=true
;   O1=false                   ;   D2=false
),                             ),
(   O1==false,                 (   D2==false,
    L1==true ->                    B2==true ->
    P1=true                        E2=true
;   P1=false                   ;   E2=false
),                             ),
(   P1==true ->                access(prime_mReset,F2),
    Q1='TooLow'                (   F2=='Off' ->
;   Q1=none                        G2=true
)                              ;   G2=false
;   Q1=I1                      ),
)                              access(prime_
)                                 mcPressure,H2),
;   Q1=none                    (   H2=='High' ->
),                                 I2=true
(   Q1==none ->                ;   I2=false
(   S='High' ->                ),
    access(prime_              (   I2==true ->
      WaterPres,R1),               J2=false
    access(prime_              ;   J2=true
        Permit,S1),            ),
(   R1<S1 ->                   (   G2==true,
    T1=true                        J2==true ->
;   T1=false                       K2=true
),                             ;   K2=false
    access(mWaterPres,         ),
        U1),                   (   E2==true,
    access('Permit',               K2==true ->
        V1),                       L2=true
(   U1<V1 ->                   ;   L2=false
    W1=true                    ),
;   W1=false                   (   L2==true ->
),                                 M2=true
(   W1==false,                 ;   M2=none
    T1==true ->                ),
    X1=true                    (   M2==none ->
;   X1=false                       access(prime
),                                     _mReset,N2),
(   X1==true ->                (   N2=='On' ->
    Y1='Permitted'                 O2=true
;   Y1=none                    ;   O2=false
)                              ),
;   Y1=none                        access(mReset,P2),
)                              (   P2=='On' ->
;   Y1=Q1                           Q2=true
)                              ;   Q2=false
;   Y1=A1                      ),
),                             (   Q2==false,
(   Y1==none ->                    O2==true ->
    access(mcPressure, Z1),        R2=true
    update(prime_              ;   R2=false
        mcPressure, Z1)        ),
;   update(prime_
```

19

```
        access(prime_
            mcPressure,S2),
    (   S2=='High' ->
        T2=true
    ;   T2=false
    ),
    (   T2==true ->
        U2=false
    ;   U2=true
    ),
    (   R2==true,
        U2==true ->
        V2=true
    ;   V2=false
    ),
    access(prime_
        mcPressure,W2),
    (   W2=='High' ->
        X2=true
    ;   X2=false
    ),
    access(mcPressure,Y2),
    (   Y2=='High' ->
        Z2=true
    ;   Z2=false
    ),
    (   Z2==false,
        X2==true ->
        A3=true
    ;   A3=false
    ),
    access(prime_
        mcPressure,B3),
    (   B3=='High' ->
        C3=true
    ;   C3=false
    ),
    (   C3==true ->
        D3=false
    ;   D3=true
    ),
    access(mcPressure,E3),
    (   E3=='High' ->
        F3=true
    ;   F3=false
    ),
    (   F3==true ->
        G3=false
    ;   G3=true
    ),
    (   G3==false,
        D3==true ->
        H3=true
    ;   H3=false
    ),
    (   A3==false,
        H3==false ->
        I3=false
    ;   I3=true
    ),
    (   V2==false,
        I3==false ->
        J3=false
    ;   J3=true

    ),
    (   J3==true ->
        K3=false
    ;   K3=none
    )
    ;   K3=M2
    ),
    (   K3==none ->
        access(tOverridden,L3),
        update(prime_
            tOverridden,L3)
    ;   update(prime_
            tOverridden,K3)
    ),
    access(mcPressure, M3),
    (   M3='TooLow' ->
        access(prime_
            tOverridden,N3),
    (   N3==true ->
        O3='Off'
    ;   O3=none
    ),
    (   O3==none ->
        access(prime_
            Overridden,P3),
    (   P3==true ->
        Q3=false ;
        Q3=true
    ),
    (   Q3==true ->
        R3='On' ;
        R3=none
    )
    ;   R3=O3
    )
    ;   R3=none
    ),
    (   R3==none ->
    (   M3='High' ->
        S3='Off'
    ;   M3='Permitted' ->
        S3='Off'
    ;   S3=none
    )
    ;   S3=R3
    ),
    (   S3==none ->
        access(cSafety_
            Injection,T3),
        update(prime_cSafety_
            Injection,T3)
    ;   update(prime_cSafety_
            Injection,S3)
    ).
```