


## SPIN Model Checker

Abhik Roychoudhury  
Department of CS  
National University of Singapore

IISc Summer Course 2007 by  
Abhik Roychoudhury

1




## The context

- A tool for modeling complex concurrent and distributed systems.
- Provides:
  - Promela, a protocol meta language
  - A model checker
  - A random simulator for system simulation
  - Promela models can be automatically generated from a safe subset of C.

IISc Summer Course 2007 by  
Abhik Roychoudhury

2




## Ideal Usage

- Write programs in C
  - Or C programs for each process in a distributed sys.
- Generate Promela Code from C automatically.
- Use the model checker of SPIN to search through the model represented by the Promela code (automatic verification).
- But ...
  - C → Promela tool relatively new.
  - Promela itself is useful for modeling protocols etc.

IISc Summer Course 2007 by  
Abhik Roychoudhury

3




## Our Usage

- Learn Promela, a low-level modeling language.
- Use it to model simple concurrent system protocols and interactions.
- Gain experience in verifying such concurrent software using the SPIN model checker.
- Gives a feel (at a small scale)
  - What are hard-to-find errors ?
  - How to find the bug in the code, once model checking has produced a counter-example ?

IISc Summer Course 2007 by  
Abhik Roychoudhury

4




## Why Promela ?

- Low-level specification language to model finite state sys.
- Models finite state concurrent processes which compute and communicate.
- Fairly extensive coverage of communication
  - Via global shared variables.
  - Via message channels
    - Synchronous communication (hand-shake)
    - Asynchronous communication (buffers)

IISc Summer Course 2007 by  
Abhik Roychoudhury

5



## Why Promela ?

- Extensive support of various control constructs for computation.
  - Assignments, Assert, If, Do
  - Ideas from guarded command languages
- Dynamic creation of processes supported.
  - Gives the flavor of a realistic multi-threaded programming language
  - Yet supported directly by a model checker !!
  - Ideal for our purposes in this course.

IISc Summer Course 2007 by  
Abhik Roychoudhury

6

## Example 0

```

byte state = 0;

proctype A()
{
  byte tmp;

  (state==0) -> tmp = state;
  tmp = tmp+1;
  state = tmp;
}

init { run A(); }

```

- **state** : Global Variable
- **tmp** : Local Variable
- **(state==0) -> tmp = state** is a guarded command (**blocked if the guard is false**).
- Only one process created.
- Final value of **state** is 1

*But SPIN allows multiple processes to be created.*

IISc Summer Course 2007 by Abhik Roychoudhury 7

## Example 1

```

byte state = 0;

proctype A()
{
  byte tmp;

  (state==0) -> tmp = state;
  tmp = tmp+1; state = tmp;
}

init { run A(); run A(); }

```

What will happen here ?

We need to define **how processes are scheduled** to determine behaviors.

IISc Summer Course 2007 by Abhik Roychoudhury 8

## Process scheduling

- All processes execute concurrently
- Interleaving semantics
  - At each time step, only one of the "active" processes will execute (**non-deterministic choice** here)
  - A process is active, if it has been created, and its "next" statement is not blocked.
  - Each statement in each process executed atomically.
  - Within the chosen process, if several statements are enabled, one of them executed non-deterministically.
    - We have not seen such an example yet !

IISc Summer Course 2007 by Abhik Roychoudhury 9

## Simulation and Verification

- At this point, please note:
  - Promela being an executable specification language, we talk of how Promela programs are executed.
  - Non-determinism broken during simulation.
  - We will mostly use Promela in a different way.
  - Given a Promela program
    - We let SPIN generate a model out of it internally.
    - This model captures ALL possible traces.
    - The model-checker traverses all possible behaviors for debugging (**diff. from execution!**)

IISc Summer Course 2007 by Abhik Roychoudhury 10

## Looking inside a process

- Data Structures
  - Basic types : int, bool, bit, byte
  - Arrays
  - Records (through typedef declarations)
- Just as in C, not much going on here !
- Check SPIN manual for details
  - <http://spinroot.com/spin/Man/Manual.html>

IISc Summer Course 2007 by Abhik Roychoudhury 11

## Statements

- Assignments
- Boolean expressions
  - **If true, then no-op else block**
- Guarded commands
  - (state == 1) -> tmp = state;
  - Guard and body evaluated separately, be careful !!
  - If you want to evaluate them together
  - atomic { (state == 1) -> tmp = state; }
  - Effect of a **test-and-set** instruction

IISc Summer Course 2007 by Abhik Roychoudhury 12

## Example 1 - Revisited

```
byte state = 0;

proctype A()
{ byte tmp;

  (state==0) -> tmp = state;
  tmp = tmp+1; state = tmp;
}

init { run A(); run A(); }
```

Final val. of state can still be 1 ??

Problem of arbitrary shared variable access by several threads.

IISc Summer Course 2007 by  
Abhik Roychoudhury

13

## Verification Example 0

```
bit flag;
byte sem;
proctype myprocess(bit i)
{ (flag != 1) -> flag = 1;
  sem = sem + 1;
  sem = sem - 1;
  flag = 0;
}
proctype observer() {
  assert( sem != 2 );
}

init {
  atomic{
    run myprocess(0);
    run myprocess(1);
    run observer();
  }
}
```

All three processes  
instantiated together

IISc Summer Course 2007 by  
Abhik Roychoudhury

14

## Issues

- Initial values of sem, flag not given
  - All possible init. values used for model checking.
- The system being verified is the asynchronous composition `myprocess(0) || myprocess(1)`
- The property is the invariant
  - $G \text{ sem} \neq 2$
- Local & global invariants can be specified inside code via assert statements.

IISc Summer Course 2007 by  
Abhik Roychoudhury

15

## More on assert

- Of the form `assert B`
  - B is a boolean expression
  - If B then no-op else abort (with error).
- Can be used inside a process (local invariants)
  - `proctype P( ... ) { x = ...; assert( x != 2 ); ... }`
- Or as a separate observer process (global invariants)
  - `proctype observer() { assert(x != 2); }`

IISc Summer Course 2007 by  
Abhik Roychoudhury

16

## Warm-up Exercise

- Try out verification example 0 in SPIN
- Try to correct the bug based on the evidence generated by the model checker.

IISc Summer Course 2007 by  
Abhik Roychoudhury

17

## Verification Example 1

```
bit flags[2];
byte sem, turn;
proctype myprocess(bit id) {
  flags[id] = 1;
  turn = 1 - id;
  flags[1-id] == 0 || turn == id;
  sem++;
  sem--;
  flags[id] = 0;
}

init() {
  atomic{
    run myprocess(0);
    run myprocess(1);
    run observer();
  }
}

proctype observer() {
  assert( sem != 2 );
}
```

IISc Summer Course 2007 by  
Abhik Roychoudhury

18

## Issues

- Can you use SPIN to prove mutual exclusion ?
  - What purpose does **turn** serve ?
- Arrays have been used in this example.
  - Flags is global, but each element is updated by only one process in the protocol
  - Not enforced by the language features.
- Processes could alternatively be started as:
  - active** proctype myprocess(...) {
  - Alternative to dynamic creation via **run** statement

IISc Summer Course 2007 by  
Abhik Roychoudhury

19

## So far ...

- Process creation and interleaving.
- Process communication via shared variables.
- Standard data structures within a process.
- Assignment, Assert, Guards.
- NOW ...
  - Guarded IF and DO statements
  - Channel Communication between processes**
  - Model checking of LTL properties

IISc Summer Course 2007 by  
Abhik Roychoudhury

20

## Non-deterministic choice

- Choice of statements within a process
  - if
  - :: condition<sub>1</sub> -> ... ; ... ; ...
  - ...
  - :: condition<sub>k</sub> -> ... ; ... ; ...
  - fi;
- If several conditions hold, select and execute any one (more behaviors for verification).
- If none hold, the statement blocks.

IISc Summer Course 2007 by  
Abhik Roychoudhury

21

## Loops

- Similar to the **if-fi** statement, we have a **do-od** statement.
- Repeat the choice selection forever.
  - Useful for modeling infinite loops pre-dominant in control software.
- Control can transfer out of the loop via a break statement in the flavor of the C language.

IISc Summer Course 2007 by  
Abhik Roychoudhury

22

## A loop which may terminate

byte count;

```
proctype counter()
{
  do
    :: count = count + 1
    :: count = count - 1
    :: (count == 0) -> break
  od;
}
```

Enumerate the reasons for non-termination in this example

IISc Summer Course 2007 by  
Abhik Roychoudhury

23

## A loop which will not terminate

```
active proctype TrafficLightController() {
  byte color = green;
  do
    :: (color == green) -> color = yellow;
    :: (color == yellow) -> color = red;
    :: (color == red) -> color = green;
  od;
}
```

IISc Summer Course 2007 by  
Abhik Roychoudhury

24

## Channels

- SPIN processes can communicate by exchanging messages across channels
- Channels are typed.
- Any channel is a FIFO buffer.
- Handshakes supported when buffer is null.
- `chan ch = [2] of bit;`
  - A buffer of length 2, each element is a bit.
- Array of channels also possible.
  - Talking to diff. processes via dedicated channels.

IISc Summer Course 2007 by  
Abhik Roychoudhury

25

## Value-passing

```
chan ch = [0] of bit;

active proctype sender()
{
    ch!1;
}

active proctype receiver()
{
    bit x;
    ch?x;
    printf("%d", x);
}
```

The value 1 is passed into local var. x via message passing.  
In this example, the message passing was via a handshake  
**! is output, ? is input**

IISc Summer Course 2007 by  
Abhik Roychoudhury

26

## Message retrieval

- `ch ? X`
  - Retrieve the earliest received (note: FIFO) message from the buffer for ch and store it into the local var. X on the receiver side.
- `ch ? 1`
  - Same as `(ch ? X; X == 1)`
- Receiving is always blocked if the corresponding channel buffer is empty.
- Similarly for sending.

IISc Summer Course 2007 by  
Abhik Roychoudhury

27

## An example with channels

```
chan name = [??] of byte;

init { atomic { run A(); run B() } }

proctype A() {
    name!124;
    name!121;
}

proctype B() {
    byte state;
    name?state;
}
```

Enumerate the behaviors when:  
?? is 0  
?? is 1  
?? is > 1

IISc Summer Course 2007 by  
Abhik Roychoudhury

28

## Another (more famous) example

```
#define p 0
#define v 1
chan sema = [0] of { bit };
proctype dijkstra_semaphore() {
    byte count = 1;
    do
        :: (count == 1) -> sema!p; count = 0
        :: (count == 0) -> sema?v; count = 1
    od
}
```

IISc Summer Course 2007 by  
Abhik Roychoudhury

29

## Another (more famous) example

```
proctype user()
{
    do
        :: sema?p; /* critical section */
        sema!v; /* non-critical section */
    od
}

init {
    run dijkstra_semaphore(); run user(); run user(); run user()
}
```

IISc Summer Course 2007 by  
Abhik Roychoudhury

30

## Example: ABP


- Alternating Bit Protocol
  - Reliable channel communication between sender and receiver.
  - Exchanging msg and ack.
  - Channels are lossy
  - Attach a bit with each msg/ack.
  - Proceed with next message if the received bit matches your expectation.
- `chan datachan = [2] of { bit };`
- `chan ackchan = [2] of { bit };`

IISc Summer Course 2007 by  
Abhik Roychoudhury

31

## ABP architecture

Implemented as SPIN processes



IISc Summer Course 2007 by  
Abhik Roychoudhury

32

## Sender

- ```

active proctype Sender()
{
  bit out, in;
  do
  :: datachan!out ->
    ackchan?in;
    if
    :: in == out -> out = 1- out;
    :: else fi
  od
}
  
```

IISc Summer Course 2007 by  
Abhik Roychoudhury

33

## Receiver

- ```

active proctype Receiver()
{
  bit in ;
  do
  :: datachan?in -> ackchan!in
  :: timeout -> ackchan!in
  od
}
  
```

IISc Summer Course 2007 by  
Abhik Roychoudhury

34

## Timeouts

- Special feature of the language
  - Time **independent** feature.
    - Do not specify a time as if you are programming.
  - True if and only if there are no executable statements in any of the currently active processes.
  - True modeling of deadlocks in concurrent systems (and the resultant recovery).

IISc Summer Course 2007 by  
Abhik Roychoudhury

35

## Handling procedures

- Processes in SPIN can be used to model and validate procedures
  - Each instantiation of a procedure (via procedure call) is modeled by spawning of a process.
- How do the called and caller processes communicate ?
  - Common channels
  - Shared Variables
- Makes SPIN useful for software verification.

IISc Summer Course 2007 by  
Abhik Roychoudhury

36

## A simple example

- The factorial function
- `int fact (int n)`
- {
- if `n <= 1` return 1;
- else return `n*fact(n-1)`;
- }

Each invocation of `fact` is a separate process. The number of processes is finite as long as the procedure call stack remains bounded.

IISc Summer Course 2007 by  
Abhik Roychoudhury

37

## Factorial in SPIN

- `proctype fact(int n; chan p)`
- { `int result;`
- if
- :: `(n <= 1) -> p!1`
- :: `(n > 1) -> chan child = [1] of {int};`
- run `fact(n-1, child);`
- `child?result;`
- `p!n*result`
- fi
- }

child serves as input channel  
p serves as output channel

IISc Summer Course 2007 by  
Abhik Roychoudhury

38

## Now initialize it as ...

- `init`
- { `int result;`
- `chan child = [1] of {int};`
- 
- run `fact(10, child);`
- `child?result;`
- `printf("result is %d\n", result);`
- `assert(result > 1024);`
- }

Computes `fact(10)` and  
verifies that it is  
greater than  $2^{10}$

IISc Summer Course 2007 by  
Abhik Roychoudhury

39

## Do not forget ...

- ... the obvious
- SPIN is a model checking tool which proceeds by finite graph search.
- Cannot be used to prove theorems like
  - For all  $n > 3$ ,  $\text{fact}(n) > 2^n$
- Proof of such theorems involve deductive machinery like mathematical induction
- Supported by theorem provers like PVS
  - To be studied later in this course.

IISc Summer Course 2007 by  
Abhik Roychoudhury

40

## Modeling Exercise

- The well-known Ackermann's function is defined as follows:
  1. If  $x = 0$  then  $A(x, y) = y + 1$
  2. If  $y = 0$  then  $A(x, y) = A(x-1, 1)$
  3. Otherwise,  $A(x, y) = A(x-1, A(x, y-1))$
- Write a Promela process "ack" to compute  $A(x, y)$ .
  - Will need to pass channels to processes.

IISc Summer Course 2007 by  
Abhik Roychoudhury

41

## Answer to Ackermann's func.

- `proctype ack(int m; int n; chan res)`
- {
- `int result;`
- `chan child = [1] of {int};`
- 
- if
- :: `(m == 0) -> res!n+1;`
- :: `(n == 0) -> run ack(m-1, 1, child);`
- `res!result;`
- :: else -> run `ack(m, n-1, child);`
- `child?result;`
- run `ack(m-1, result, child);`
- `child?result;`
- `res!result;`
- fi
- }

- `init {`
- `int result;`
- `chan res = [1] of {int};`
- run `ack(4, 4, res);`
- `res?result;`
- `printf("Result is %d\n", result);`
- }

IISc Summer Course 2007 by  
Abhik Roychoudhury

42

## More Modeling Exercises

- Use SPIN to prove mutual exclusion of the semaphore encoding.
- Enough of modeling, let us do some verification.
- Features of PROMELA relevant to verification
  - End, Progress, Accept labels

IISc Summer Course 2007 by  
Abhik Roychoudhury

43

## Part II: Verification using SPIN

Abhik Roychoudhury  
Department of CS  
National University of Singapore

IISc Summer Course 2007 by  
Abhik Roychoudhury

44

## Quotable Quotes

- "I have been fishing all day, I have found a number of fish since the morning, I cannot find any more now, I am pretty sure, there aren't any left!"
  - Folklore
    - Taken from Antonia Bertolino's slides on testing
- Bug finding techniques will ensure worse coverage than fishing in a small pond.

IISc Summer Course 2007 by  
Abhik Roychoudhury

45

## Quotable Quotes

- "If I had eight hours to chop down a tree, I would spend six hours sharpening my axe."
  - U.S. President Abraham Lincoln
    - 1809 – 1865
- Time investment in building verifiers is time well-spent!

IISc Summer Course 2007 by  
Abhik Roychoudhury

46

## Execution engine

- Select an enabled transition of any thread, and execute it.
- A transition corresponds to one statement in a thread.
  - Handshakes must be executed together.
    - `chan x = [0] of {...};`
    - `x!1`            `// x?data`

IISc Summer Course 2007 by  
Abhik Roychoudhury

47

## Execution engine

```
while ( (E = executable(s)) != {} )
  for some (p,t) ∈ E
  {
    s' = apply(t.effect, s); /* execute the chosen statement */
    if (handshake == 0)
    {
      s = s';
      p.curstate = t.target;
    }
    else{ ...
  }
```

IISc Summer Course 2007 by  
Abhik Roychoudhury

48



## Execution engine

```

/* try to complete the handshake */
E' = executable(s); /* E' = {}  $\Rightarrow$  s unchanged */
for some (p', t')  $\in$  E'
{
  s = apply(t'.effect, s);
  p.curstate = t.target;
  p'.curstate = t'.target;
}
handshake = 0
} /* else */
} /* for some (p, t)  $\in$  E */
} /* while ((E = executable(s)) ... */
while (stutter) { s = s }

```

IISc Summer Course 2007 by  
Abhik Roychoudhury

49

## Specifying Properties in SPIN

- Invariants
  - Local: via assert statement insertion
  - Global: assert statement in a monitor process
- Deadlocks
- Bad Cycles
- Arbitrary Temporal Properties
  - SPIN is a LTL model checker.

IISc Summer Course 2007 by  
Abhik Roychoudhury

50

## Deadlocks

- When all processes are blocked.
- Exhibited by
  - Finite execution traces where all processes instantiated have not terminated and are blocked
- But all processes in a PROMELA program may not be meant to terminate !
  - Our Traffic Light Controller example
- Specify legal end-states of the processes
  - And modify the detection of deadlock as ...

IISc Summer Course 2007 by  
Abhik Roychoudhury

51

## Deadlocks

- Exhibited by –
  - Finite execution traces where all instantiated processes have not terminated and not reached a legal end-state, and are blocked.
- Semaphore example
  - proctype semaphore()
  - { byte count = 1;
  - end: do
  - :: (count == 1) -> sema!p; count = 0
  - :: (count == 0) -> sema?v; count = 1
  - od }

IISc Summer Course 2007 by  
Abhik Roychoudhury

52

## Deadlock detection

- As in prev. slide (any finite trace satisfying ...)
- We have marked the beginning of the infinite loop as a legal end-state of the semaphore process.
- The semaphore process is simply waiting in the loop for user requests, hence cannot contribute to a deadlock.
- There can be multiple end-states in a process
  - Check SPIN manual on how to mark them.

IISc Summer Course 2007 by  
Abhik Roychoudhury


53

## No-progress Cycles

- An infinite loop where processes execute actions, but no “progress” is achieved.
- Example:
  - A communication protocol where the parties keep on exchanging control signals, but no data is actually communicated.
- Need to clarify what is “progress”
  - By inserting progress labels in the Promela model.

IISc Summer Course 2007 by  
Abhik Roychoudhury

54



## No progress cycles

```

proctype dijkstra()
{
  byte count = 1;
  end: do
    :: (count == 1) ->
    progress: sema!p; count = 0
    :: (count == 0) -> sema?v; count = 1
  od
}

```

Verifies that at least one process enters the critical section.

IISc Summer Course 2007 by  
Abhik Roychoudhury

56

## Correctness claims

- Progress labels
  - Any infinite execution cycle contains at least one progress label.
  - No progress cycles are cycles without any progress label
- Acceptance labels
  - No execution trace passes through an accept label infinitely often
  - Model Checking reports an acceptance cycle (if any)
  - Acceptance cycles are cycles with at least one acceptance state.

IISc Summer Course 2007 by  
Abhik Roychoudhury

57

## Acceptance cycles

- A cycle which goes through an "acceptance" state infinitely often
  - A "bad" cycle if the acceptance state is supposed to occur only finitely many times.
  - An acceptance state could mark the state reached after some initialization activity in a protocol.
  - Accept. Cycle  $\Rightarrow$  System unintentionally getting reset!
- Can mark acceptance states by "accept" labels in Promela code
  - Labels can be marked by user
  - Accept labels can be automatically generated from user-provided LTL properties to support LTL verification (later !)

IISc Summer Course 2007 by  
Abhik Roychoudhury

58

## Acceptance cycles

```

proctype semaphore()
{
  byte count = 1;
  end: do
    :: (count == 1) ->
    progress: sema!p; count = 0
    :: (count == 0) ->
    accept:  sema?v; count = 1
  od
}

```

The acceptance label makes it impossible to loop through P and V operations of the semaphore - this property is false incidentally.

IISc Summer Course 2007 by  
Abhik Roychoudhury

59

## Model Checking

- $(P1 \parallel P2 \parallel P3) \models \varphi$ 
  - P1, P2, P3 are Promela processes
  - $\varphi$  is a LTL formula
- Construct a state machine via
  - M, asynchronous composition of processes P1, P2, P3
  - $M(\neg\varphi)$ , representing  $\neg\varphi$
- Show that "language" of  $M \times M(\neg\varphi)$  is empty

IISc Summer Course 2007 by  
Abhik Roychoudhury

60

## Model Checking

- A given LTL property (rather its negation) is internally represented as an automata.
- This property automata is **synchronously** composed with the global system automata.
- We then show that the traces accepted by the composition of the system and property automata is empty.
- But the traces are potentially infinite ...
  - Finite state automata over infinite inputs


IISc Summer Course 2007 by  
Abhik Roychoudhury

61

## Buchi automata

Mark certain states as acceptance states as usual.

Accept an infinite string if (at least one of) its runs through the automata visits (one or more of) the acceptance states infinitely often. OTHERWISE reject the string.



Accepts infinite strings with finitely many occurrences of p

IISc Summer Course 2007 by  
Abhik Roychoudhury

62

## Buchi automata

- Like conventional finite-state automata
  - $A = (S, \Sigma, I, \rightarrow, F)$ 
    - $S$ , set of states
    - $\Sigma$ , a finite alphabet
    - $I \subseteq S$ , set of initial states
    - $\rightarrow \subseteq S \times \Sigma \times S$ , transition relation
    - $F \subseteq S$ , set of final states
  - Notion of acceptance is different

IISc Summer Course 2007 by  
Abhik Roychoudhury

63

## Buchi automata

- Run  $r$  of a string  $\sigma \in \Sigma^\omega$ 
  - Sequence of states of  $A$  obtained by running  $\sigma$  from an initial state of automata  $A$
  - $r[0] \in I$  and, for all  $i \geq 0$ ,  $r[i] \xrightarrow{\sigma[i]} r[i+1]$
- Given a run  $r$ ,
  - $\text{inf}(r)$  = set of states appearing infinitely often in  $r$
  - These are the states that are visited infinitely often on running the infinite string  $\sigma$
- Language of the automata (notion of acceptance)
  - $L(A) = \{\sigma \mid \sigma \in \Sigma^\omega \text{ and } \sigma \text{ has a run } r \text{ s.t. } \text{inf}(r) \cap F \neq \emptyset\}$

IISc Summer Course 2007 by  
Abhik Roychoudhury

64


## Buchi automata

- Conventional finite state automata over finite strings
  - String accepted if it ends in a final state
- Buchi automata over infinite strings
  - String accepted if it visits at least one final state infinitely often.
    - We need to deal with infinite strings since the system execution traces are infinite.

IISc Summer Course 2007 by  
Abhik Roychoudhury

65

## LTL properties



Corresponds to the negation of the LTL property


$GF\ p$  ( assuming  $\Sigma = \{p, q\}$  )

If the user seeks to verify  $GF\ p$ , SPIN generates Promela code for the negation of the property which will internally construct such an automata.

IISc Summer Course 2007 by  
Abhik Roychoudhury

66

## A more complex property



Represents negation of the LTL property  
 $G ( p \Rightarrow (p \cup q) )$

Internally generated by SPIN when the user wants to verify the LTL property.

Acceptance states correspond to **accept** labels discussed earlier.

IISc Summer Course 2007 by  
 Abhik Roychoudhury

67

## Verif. via Acceptance cycles

- Given the Buchi automata for the negation of LTL property (and its acceptance states)
  - SPIN computes a synchronous product of this with the global transition system
  - The property automata should always make a move with the system automata
  - The language of the product automata is non-empty iff it makes the property automata move in a cycle containing acceptance states.
  - Verification achieved by nested depth first search to find such **acceptance cycles**.

IISc Summer Course 2007 by  
 Abhik Roychoudhury

68

## No-progress cycles

- Absence of no-progress cycles described in LTL as
  - GF progress**
- Negation of the property is
  - FG no\_progress**
    - where no\_progress is an atomic proposition which is true in any state where the control location is not marked as **progress**
- We can compose the program model M with the automata derived from FG no\_progress and perform model checking by detecting acceptance cycles.

IISc Summer Course 2007 by  
 Abhik Roychoudhury

69

## Important Clarification

- SPIN supports model-checking of arbitrary LTL properties by
  - Converting negation of property
  - Converting negated property to Buchi automata
  - Constructing synchronous product of design's transition system and Buchi automata of negated property
  - Defining accepting states of the Buchi automata to accept labels of the product automata, and
  - Searching for acceptance cycles in the product automata.
- Thus, accept labels are generated automatically from LTL property, and are not directly given by user.

IISc Summer Course 2007 by  
 Abhik Roychoudhury


70

## Finding acceptance cycles

- We have reduced LTL model checking to finding acceptance cycles
- How to find acceptance cycles ?
  - One possibility is by SCC detection
    - 1. Compute **strongly connected components** of the product graph (DFS)
    - 2. Check whether any SCC contains an acceptance state; if yes, an acceptance cycle exists.
  - But ...

IISc Summer Course 2007 by  
 Abhik Roychoudhury

71



IISc Summer Course 2007 by  
 Abhik Roychoudhury

72

## SPIN model checking

- SPIN does not use SCC detection for detecting acceptance cycles (and hence model checking)
- The nested DFS algorithm used in SPIN is more space efficient in practice.
  - SCC detection maintains **two integer numbers per node**. (*dfs* and *lowlink* numbers)
  - Nested DFS maintains only one integer.
    - This optimization is important due to the huge size of the product graph being traversed on-the-fly by model checker.

IISc Summer Course 2007 by  
Abhik Roychoudhury


73

## Nested DFS in SPIN

- Find acceptance states reachable from initial states (DFS).
- Find all such acceptance states which are reachable from itself (DFS).
- Counter-example evidence (if any) obtained by simply concatenating the two DFS stacks.

IISc Summer Course 2007 by  
Abhik Roychoudhury

74



IISc Summer Course 2007 by  
Abhik Roychoudhury

75

## Standard DFS on $M \times M(\varphi)$

- procedure *dfs*(*s*)
  - push *s* to Stack
  - add {*s*} to States
  - for each transition  $s \rightarrow s'$  do
    - if  $s' \notin \text{States}$  then *dfs*(*s'*)
  - endfor
  - pop *s* from Stack
- end

IISc Summer Course 2007 by  
Abhik Roychoudhury

76

## Nested DFS– Step 1

- procedure *dfs1*(*s*)
  - push *s* to Stack1
  - add {*s*} to States1
  - if accepting(*s*) then
    - States2 := empty; seed := *s*; *dfs2*(*s*)
  - endif
  - for each transition  $s \rightarrow s'$  do
    - if  $s' \notin \text{States1}$  then *dfs1*(*s'*)
  - endfor
  - pop *s* from Stack1
- end

IISc Summer Course 2007 by  
Abhik Roychoudhury

77

## Nested DFS – Step 2


- procedure *dfs2*(*s*)
  - push *s* to Stack2
  - add {*s*} to States2
  - for each transition  $s \rightarrow s'$  do
    - if  $s' = \text{seed}$  then **report acceptance cycle**
    - else if  $s' \notin \text{States2}$  then *dfs2*(*s'*)
  - endif
  - endfor
  - pop *s* from Stack2
- end

IISc Summer Course 2007 by  
Abhik Roychoudhury

78

## More on nested DFS

- Space efficient compared to SCC detection
  - Only one id maintained (the DFS number) for each state being visited.
- Time inefficient
  - Different invocations of DFS2 may search the same portion of state space.
  - But DFS2 always starts from scratch.
  - Exercise: how to re-use results of past DFS2 invocations ?



Abhik Roychoudhury

79

## Some Common Questions

- How does the product of the system and property automata work ?
  - How is the interaction between system and property automata achieved ?
- Can we specify LTL properties directly ?
  - Yes, you can do so in SPIN.
- Can we model/verify pgms with procedures
  - Yes.

IISc Summer Course 2007 by  
Abhik Roychoudhury

80

## Connect System and Property

- System model
  - int x = 100;
  - active proctype A()
    - { do
    - :: x % 2 -> x = 3\*x+1
    - od
    - }
  - active proctype B()
    - { do
    - :: !(x%2) -> x = x/2
    - od
    - }
- Property
  - GF (x = 1)
  - Insert into code
    - #define q (x == 1)
  - Now try to verify GF q

IISc Summer Course 2007 by  
Abhik Roychoudhury

81

## Enforcing fairness

All LTL counter-examples are acceptance cycles.

No point in reporting acceptance cycles which arise out of unfair scheduling.

If there are K active concurrent processes in the Promela model being verified, all of these processes should have transitions within the acceptance cycle found.

--- Weak fairness requirement

How to enforce this requirement?


Linear blow-up in state space.

Create (K+2) copies of  $(P_1 || \dots || P_K) \times M(\neg\varphi)$

$\varphi$  is the property being verified  
IISc Summer Course 2007 by  
Abhik Roychoudhury

82

## Enforcing fairness




Also, if state s in  $i^{\text{th}}$  copy has no outgoing transition by  $P_i$  add the transition

s in  $i^{\text{th}}$  copy  $\rightarrow$  s in  $(i+1)^{\text{th}}$  copy

IISc Summer Course 2007 by  
Abhik Roychoudhury

83

## Enforcing fairness



IISc Summer Course 2007 by  
Abhik Roychoudhury

84

## Let us finish with a real-life situation

- July 4, 1997
  - NASA's Pathfinder landed on Mars.
  - Tremendous engineering feat.
  - Hard to design the control software with concurrency and priority driven scheduling of threads.
  - The SpaceRover would lose contact with earth in unpredictable moments.

IISc Summer Course 2007 by  
Abhik Roychoudhury

85

## The Mars Pathfinder problem

"But a few days into the mission, not long after Pathfinder started gathering meteorological data, the spacecraft began experiencing total system resets, each resulting in losses of data. The press reported these failures in terms such as "software glitches" and "the computer was trying to do too many things at once". ...



IISc Summer Course 2007 by  
Abhik Roychoudhury

86

## Essence of the problem in SPIN

```

mtype = { free, busy, idle, waiting, running };
mtype H = idle; mtype L = idle; mtype mutex = free;

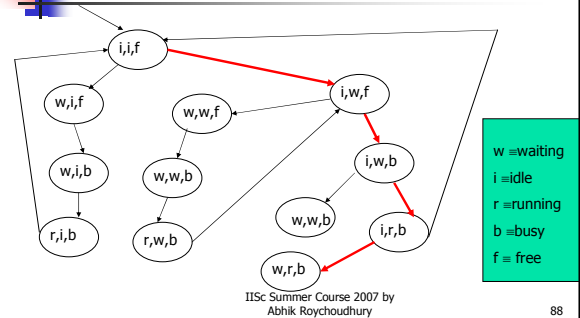
active proctype high();
{end: do
    :: H = waiting;
    atomic { mutex == free ->
        mutex = busy };
    H = running;
    atomic{ H=idle; mutex=free }
od
}

active proctype low() provided (H == idle)
{ end: do
    :: L = waiting;
    atomic{ mutex== free->
        mutex = busy};
    L = running;
    atomic{ L=idle; mutex = free }
od
}
    
```

IISc Summer Course 2007 by  
Abhik Roychoudhury

87

## State Space Graph



IISc Summer Course 2007 by  
Abhik Roychoudhury

88

## Source of deadlock

- Counterexample
  - Low priority thread acquires lock
  - High priority thread starts
  - Low priority process cannot be scheduled
  - High priority thread blocked on lock
- Actual error was a bit more complex with three threads of three different priorities
  - Timer went off with such a deadlock resulting in a system reset and loss of transmitted data.

IISc Summer Course 2007 by  
Abhik Roychoudhury

89

## The actual error

- "VxWorks provides preemptive priority scheduling of threads. Tasks on the Pathfinder spacecraft were executed as threads with priorities that were assigned in the usual manner reflecting the relative urgency of these tasks."
- "Pathfinder contained an "information bus", which you can think of as a shared memory area used for passing information between different components of the spacecraft."
  - A bus management task ran frequently with high priority to move certain kinds of data in and out of the information bus. Access to the bus was synchronized with mutual exclusion locks (mutexes)."

IISc Summer Course 2007 by  
Abhik Roychoudhury

90

## The actual error

- The meteorological data gathering task ran as an infrequent, low priority thread, ... When publishing its data, it would acquire a mutex, do writes to the bus, and release the mutex. ...
- The spacecraft also contained a communications task that ran with medium priority."

High priority: retrieval of data from shared memory  
Medium priority: communications task  
Low priority: thread collecting meteorological data

IISc Summer Course 2007 by  
Abhik Roychoudhury

91

## The actual error

- "Most of the time this combination worked fine. However, very infrequently it was possible for an interrupt to occur that caused the (medium priority) communications task to be scheduled during the short interval while the (high priority) information bus thread was blocked waiting for the (low priority) meteorological data thread. In this case, the long-running communications task, having higher priority than the meteorological task, would prevent it from running, consequently preventing the blocked information bus task from running. After some time had passed, a watchdog timer would go off, notice that the data bus task had not been executed for some time, conclude that something had gone drastically wrong, and initiate a total system reset. This scenario is a classic case of priority inversion."

IISc Summer Course 2007 by  
Abhik Roychoudhury

92

## Exercise

- Try to formalize the Mars Pathfinder problem (with 3 processes) in SPIN, for the informal discussion given here.
- Can you locate the error scenario mentioned?

IISc Summer Course 2007 by  
Abhik Roychoudhury

93

## Readings

- <http://spinroot.com/spin/Man/Manual.html>
  - SPIN manual
- The model checker SPIN (Holzmann)
  - IEEE transactions on software engineering, 23(5), 1997.
- <http://spinroot.com/spin/Doc/SpinTutorial.pdf>
  - SPIN beginner's tutorial (Theo Ruys)
- Summer school Lecture notes on Software MC
  - See Section 2, Posted under lesson plan in course web-page
- The SPIN model checker: primer and reference manual, by Holzmann (mostly chapters 2,3,7,8)

IISc Summer Course 2007 by  
Abhik Roychoudhury

94